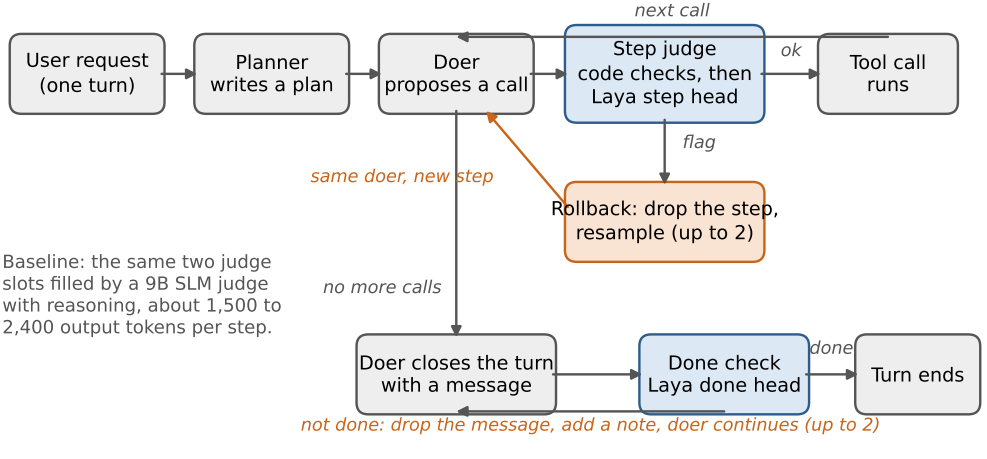
next call
User request (one turn)
Planner writes a plan
Doer proposes a call
Step judge code checks, then Laya step head
ok
Tool call runs
flag
same doer, new step
Rollback: drop the step, resample (up to 2)
Baseline: the same two judge slots filled by a 9B SLM judge with reasoning, about 1,500 to 2,400 output tokens per step.
no more calls
Doer closes the turn with a message
Done check Laya done head
done
Turn ends
not done: drop the message, add a note, doer continues (up to 2)