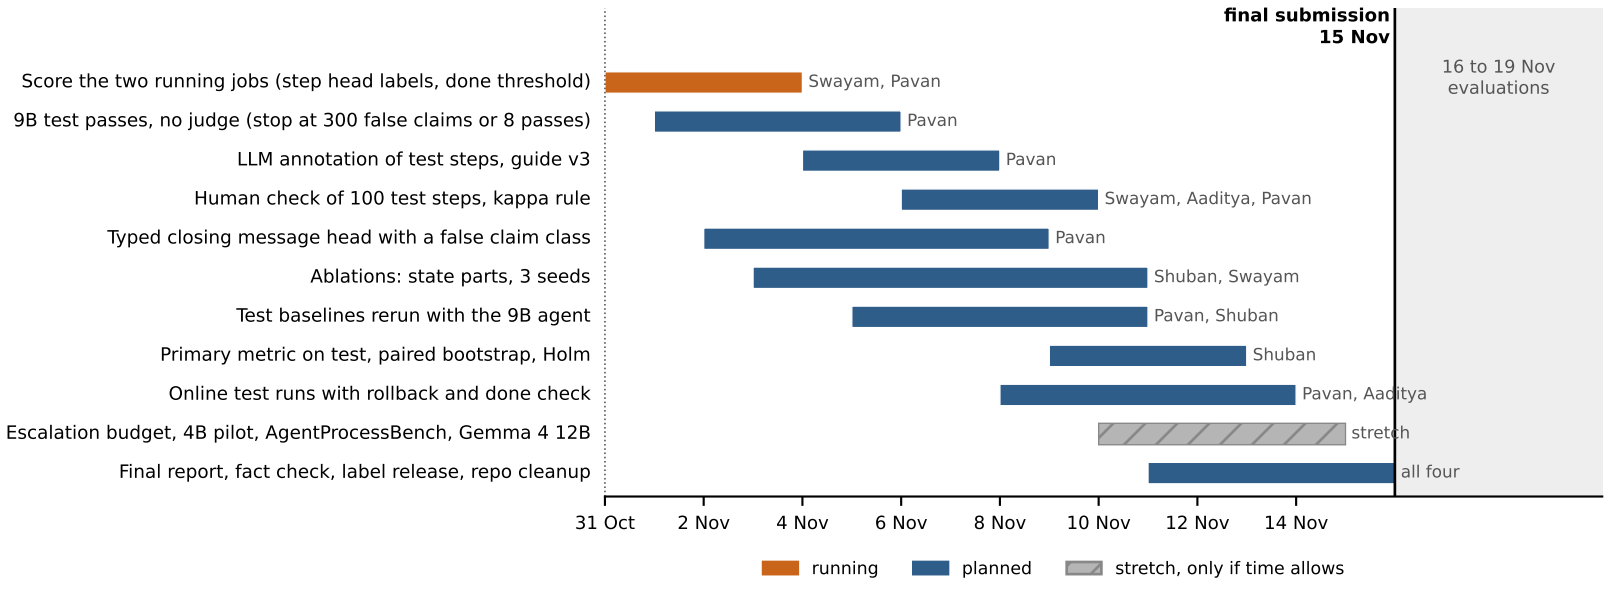
final submission
15 Nov
16 to 19 Nov
evaluations
Score the two running jobs (step head labels, done threshold)
Swayam, Pavan
9B test passes, no judge (stop at 300 false claims or 8 passes)
Pavan
LLM annotation of test steps, guide v3
Pavan
Human check of 100 test steps, kappa rule
Swayam, Aaditya, Pavan
Typed closing message head with a false claim class
Pavan
Ablations: state parts, 3 seeds
Shuban, Swayam
Test baselines rerun with the 9B agent
Pavan, Shuban
Primary metric on test, paired bootstrap, Holm
Shuban
Online test runs with rollback and done check
Pavan, Aaditya
Escalation budget, 4B pilot, AgentProcessBench, Gemma 4 12B
stretch
Final report, fact check, label release, repo cleanup
all four
31 Oct
2 Nov
4 Nov
6 Nov
8 Nov
10 Nov
12 Nov
14 Nov
running
planned
stretch, only if time allows