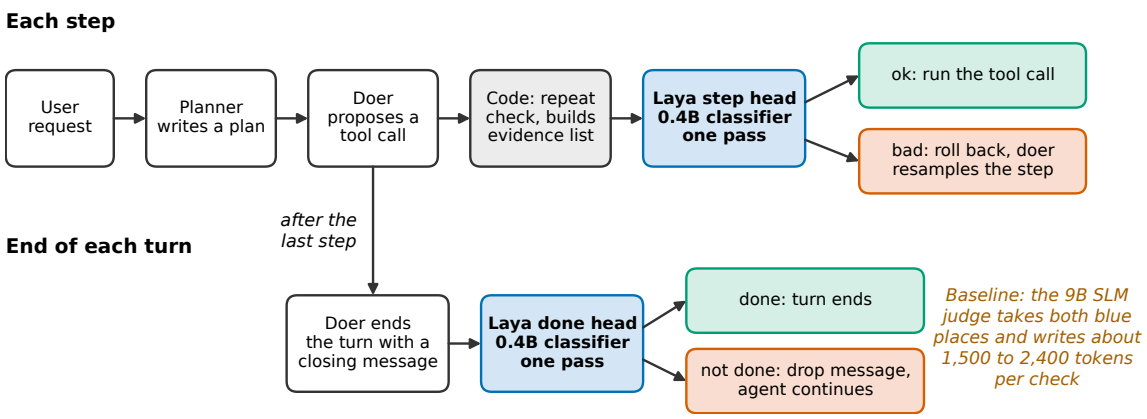
Each step
User request
Planner writes a plan
Doer proposes a tool call
Code: repeat check, builds evidence list
Laya step head 0.4B classifier one pass
ok: run the tool call
bad: roll back, doer resamples the step
after the last step
End of each turn
Doer ends the turn with a closing message
Laya done head 0.4B classifier one pass
done: turn ends
not done: drop message, agent continues
Baseline: the 9B SLM judge takes both blue places and writes about 1,500 to 2,400 tokens per check